

Great Ideas in Computer Architecture

C: Memory Management and Usage

Instructor: Shreyas Chand



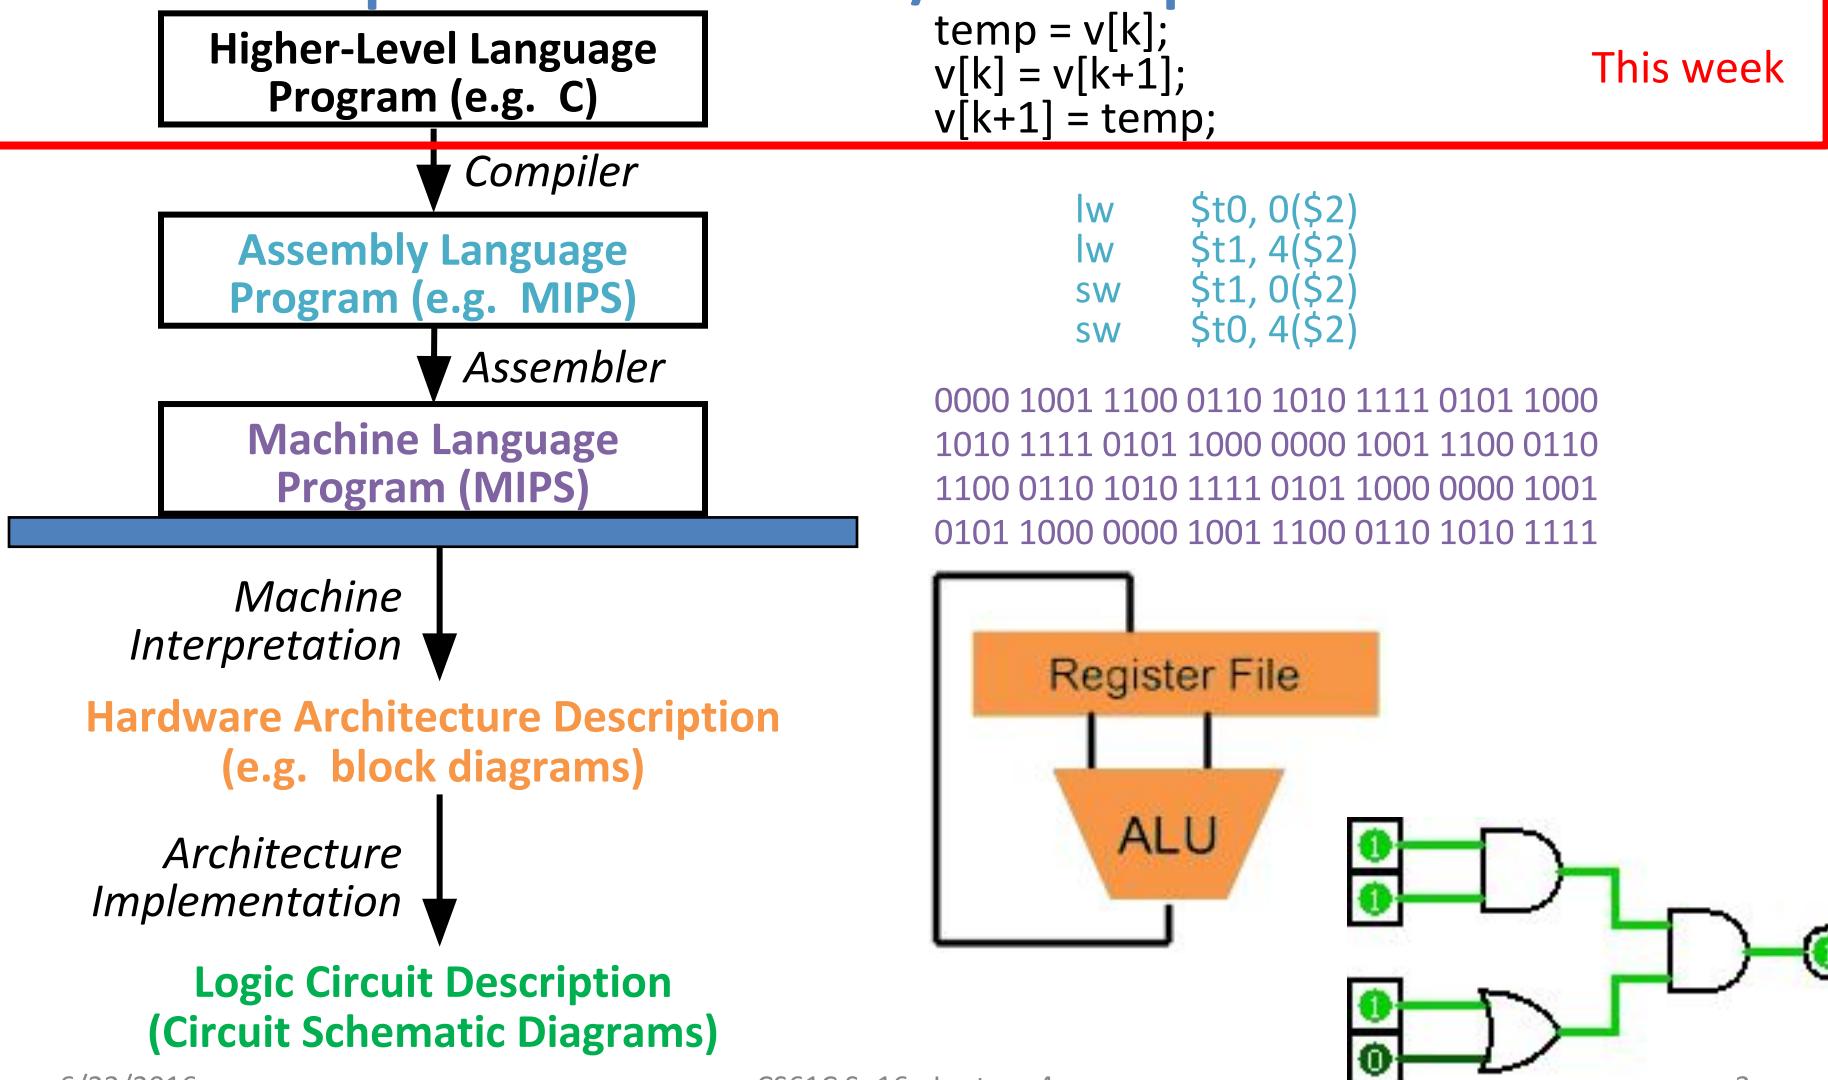
China Sunway TaihuLight officially became the fastest supercomputer in the world

- Theoretical peak performance of 125 petaflops, 10,649,600 cores, and 1.31 petabytes of primary memory.
- TaihuLight's stewards ... [are] putting all that power toward advanced manufacturing, Earth-system modeling and weather forecasting, life science, and big data analytics.
- China has more supercomputers than any other nation; Sunway is five times faster than anything in US; uses all Chinese manufactured hardware

Review

- Pointers and arrays are very similar
- Strings are just char arrays with a null terminator
- Pointer arithmetic moves the pointer by the size of the thing it's pointing to
- C accommodates for pointing to structs and pointers
- Pointers are the source of many C bugs!

Great Idea #1: Levels of Representation/Interpretation

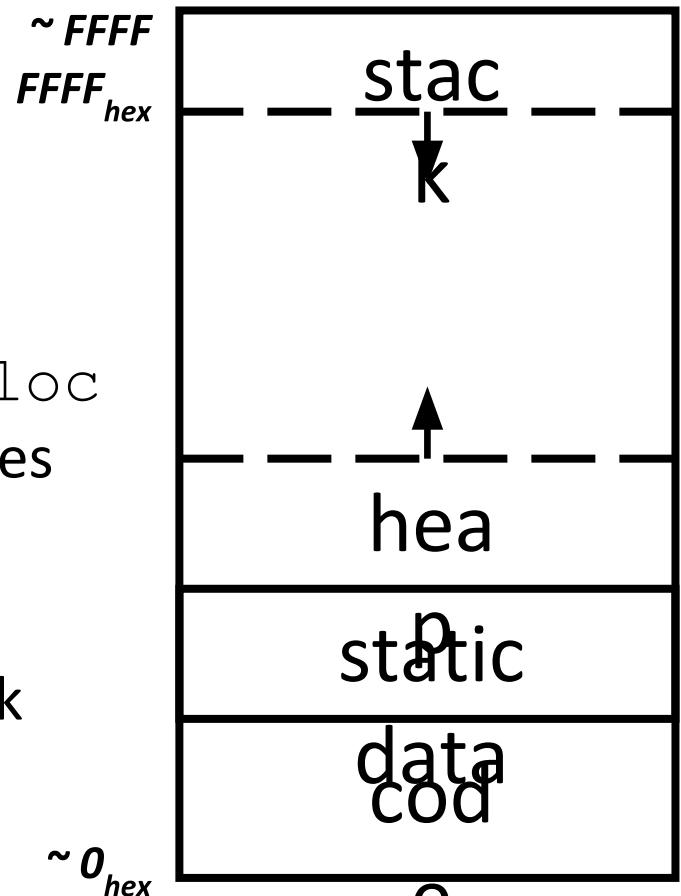


Agenda

- C Memory Layout
 - Stack, Static Data, and Code
- Administrivia
- Dynamic Memory Allocation
 - Heap
- Common Memory Problems
- C Wrap-up: Linked List Example

C Memory Layout

- Program's *address space* contains 4 regions:
 - **Stack**: local variables, grows downward
 - **Heap**: space requested via `malloc()` and used with pointers; resizes dynamically, grows upward
 - **Static Data**: global and static variables, does not grow or shrink
 - **Code**: loaded when program starts, does not change



Where Do the Variables Go?

- Declared outside a function:

Static Data

- Declared inside a function:

Stack

- main() is a function
 - Freed when function returns

- Dynamically allocated:

Heap

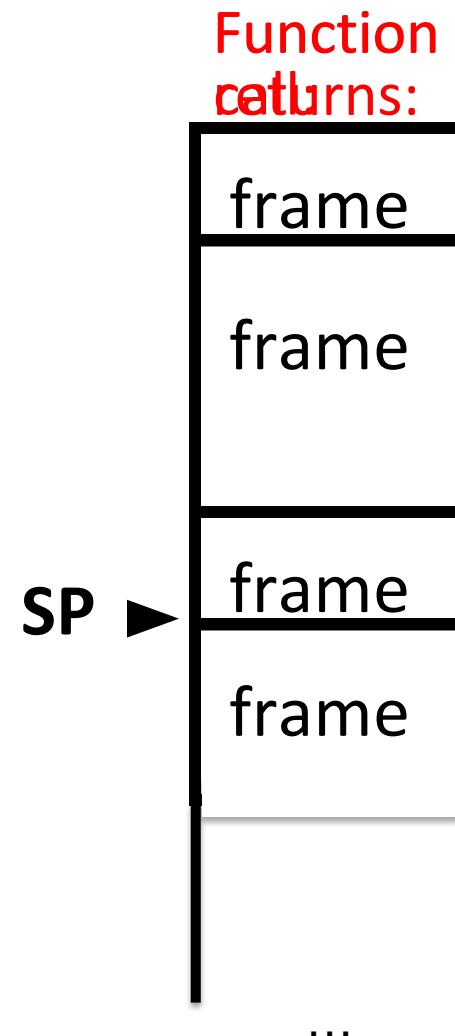
- i.e. malloc (we will cover this shortly)

```
#include <stdio.h>
int varGlobal;

int main() {
    int varLocal;
    int *varDyn =
        malloc(sizeof(int));
}
```

The Stack

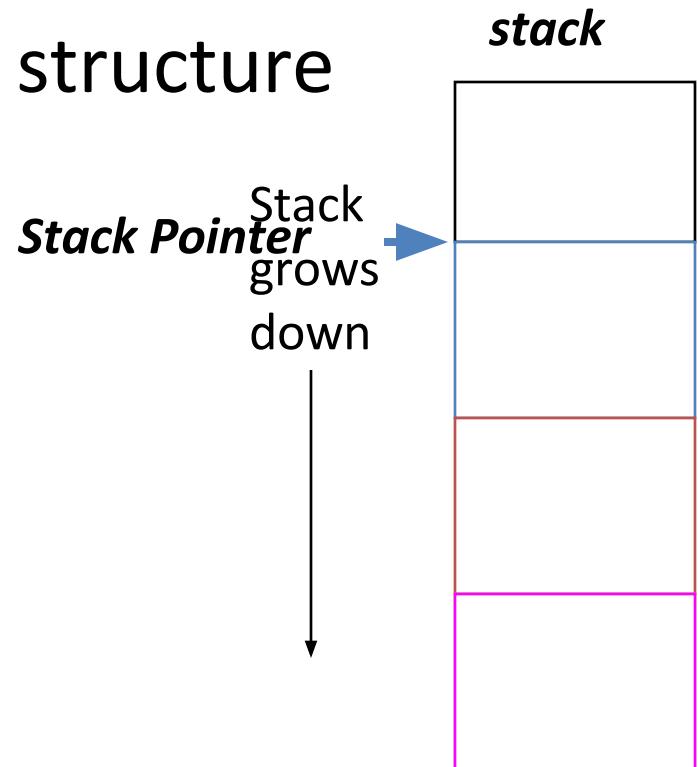
- Each stack frame is a contiguous block of memory holding the local variables of a single procedure
- A stack frame includes:
 - Location of caller function
 - Function arguments
 - Space for local variables
- Stack pointer (SP) tells where lowest (current) stack frame is
- When procedure ends, stack pointer is moved back (but data remains (**garbage!**)); frees memory for future stack frames;



The Stack

- Last In, First Out (LIFO) data structure

- int main() {
 a(0);
 return 1; }
- void a(int m) {
 b(1); }
- void b(int n) {
 c(2);
 d(4); }
- void c(int o) {
 printf("c"); }
- void d(int p) {
 printf("d"); }



Stack Misuse Example

```
int *getPtr() {  
    int y;           ←  
    y = 3;  
    return &y;  
};
```

**Never return pointers to
local variable from functions**

Your compiler will warn you about
this

```
int main () {  
    int *stackAddr, content;  
    → stackAddr = getPtr ();  
    → content = *stackAddr;  
    → printf ("%d", content); /* 3 */  
    content = *stackAddr;  
    printf ("%d", content); /* 0 */  
};
```

– don't ignore such warnings!
*overwrites
stack frame*

Static Data

- Place for variables that persist
 - Data not subject to comings and goings like function calls
 - Examples: String literals, global variables
- Size does not change, but its data can

Code

- Copy of your code goes here
 - C code becomes data too!
- Does not change

Question: Which statement below is FALSE?

All statements assume each variable exists.

```
void funcA() {int x; printf("A");}  
void funcB() {  
    int y;  
    printf("B");  
    funcA();  
}  
void main() {char *s = "s"; funcB();}
```

- (A) *x* is at a lower address than *y*
- (B) *x* and *y* are in adjacent frames
- (C) *x* is at a lower address than **s*
- (D) *y* is in the 2nd frame from the top of the Stack

Question: Which statement below is FALSE?

All statements assume each variable exists.

```

void funcA() {int x; printf("A");}
void funcB() {
    int y;
    printf("B");
    funcA();
}
void main() {char *s = "s"; funcB();}
```

This is a string literal, and thus stored in STATIC DATA.

(A) *x* is at a lower address than *y*

Note: We're talking about **s*, not *s*, i.e. the location where *s* points!

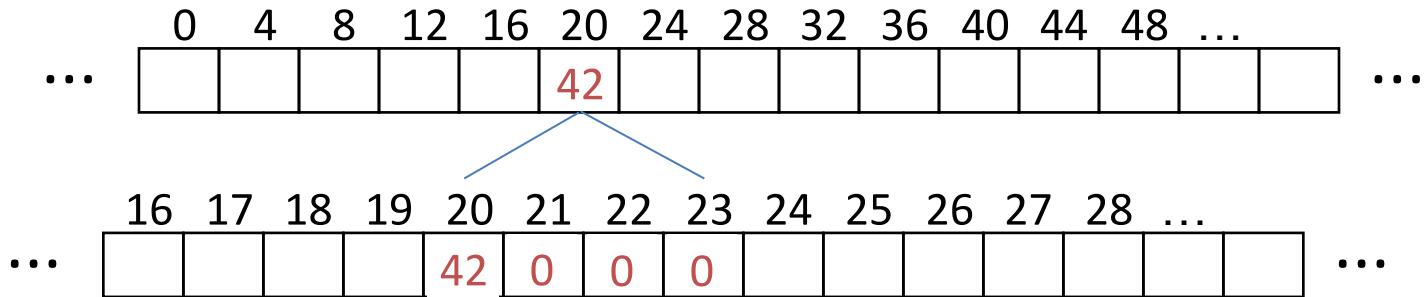
(B) *x* and *y* are in adjacent frames

(C) *x* is at a lower address than **s*

(D) *y* is in the 2nd frame from the top of the Stack

Endianness

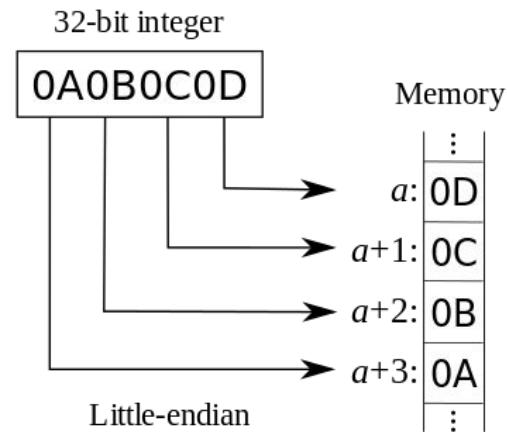
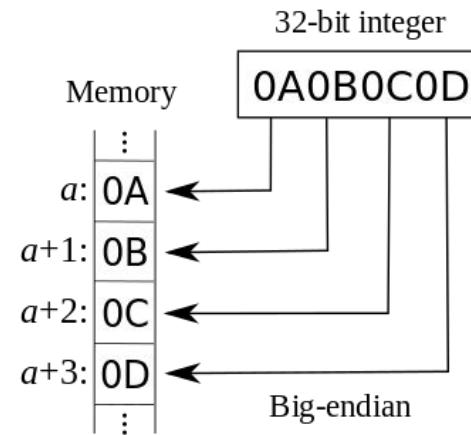
- In what order are the bytes within a data type stored in memory?



- Big Endian:
 - Descending numerical significance with ascending memory addresses
- Little Endian
 - Ascending numerical significance with ascending memory addresses

Endianness

- Big Endian:
 - Descending numerical significance with ascending memory addresses
- Little Endian
 - Ascending numerical significance with ascending memory addresses



Source: <https://en.wikipedia.org/wiki/Endianness>

Agenda

- C Memory Layout
 - Stack, Static Data, and Code
- **Administrivia**
- Dynamic Memory Allocation
 - Heap
- Common Memory Problems
- C Wrap-up: Linked List Example

Administrivia

- Meet your fellow classmates! Form study groups and get your questions answered
 - Utilize Piazza, labs, discussions, and OHs
- End of the first week!
 - HW0 and mini-bio due Sunday
 - HW1 released. Proj1 coming!
 - TA and Instructor response time increase drastically (slower) over the weekend
- Suggestion for weekend:
 - Finish HW, catch up on K&R, practice C to prepare for project!

Agenda

- C Memory Layout
 - Stack, Static Data, and Code
- Administrivia
- **Dynamic Memory Allocation**
 - Heap
- Common Memory Problems
- C Wrap-up: Linked List Example

Dynamic Memory Allocation

- Sometimes you don't know how much memory you need beforehand
 - e.g. input files, user input
- Dynamically allocated memory goes on the **Heap** – more permanent than Stack
- Need as much space as possible without interfering with Stack
 - Start at opposite end and grow towards Stack

Allocating Memory

- 3 functions for requesting memory:

`malloc()`, `calloc()`, and `realloc()`

- http://en.wikipedia.org/wiki/C_dynamic_memory_allocation#Overview_of_functions

- **malloc(*n*)**

- Allocates a continuous block of ***n bytes*** of uninitialized memory (contains garbage!)
 - Returns a pointer to the beginning of the allocated block; NULL indicates failed request (check for this!)
 - Different blocks not necessarily adjacent

Using malloc()

- Almost always used for arrays or structs
- Good practice to use `sizeof()` and typecasting

```
int *p = (int *) malloc(n*sizeof(int));
```

- `sizeof()` makes code more portable
- `malloc()` returns `void *`, typecast will help you catch coding errors when pointer types don't match
- Can use array or pointer syntax to access
- Make sure you don't lose the original address
 - `p++` is a **BAD IDEA**; use a separate pointer

Releasing Memory

- Release memory on the Heap using `free()`
 - Memory is limited, release when done
- **free(p)**
 - Pass it pointer `p` to beginning of allocated block; releases the whole block
 - `p` must be the address *originally* returned by `m/c/realloc()`, otherwise throws system exception
 - Don't call `free()` on a block that has already been released or on `NULL`

Dynamic Memory Example

- Need `#include <stdlib.h>`

```
typedef struct {  
    int x;  
    int y;  
} point;  
  
point *rect; /* opposite corners = rectangle */  
...  
if( !(rect=(point *) malloc(2*sizeof(point))))  
{  
    printf("\nOut of memory!\n");  
    exit(1);  
}  
    ← Do NOT change rect during this time!!!  
...  
free(rect);
```

Check for
returned NULL

Question: Want output: $a[] = \{0,1,2\}$ with no errors.

Which lines do we need to change?

```
1 #define N 3
2 int *makeArray(int n) {
3     int *ar;
4     ar = (int *) malloc(n);
5     return ar;
6 }
7 void main() {
8     int i, *a = makeArray(N);
9     for(i=0; i<N; i++)
10         *a++ = i;
11     printf("a[] =
12         { %i, %i, %i } ", a[0], a[1], a[2]);
13     free(a);
14 }
```

- (A) 4, 12
- (B) 5, 12
- (C) 4, 10
- (D) 5, 10

Question: Want output: $a[] = \{0,1,2\}$ with no errors.

Which lines do we need to change?

```

1 #define N 3
2 int *makeArray(int n) {
3     int *ar;
4     ar = (int *) malloc(n * sizeof(int));
5     return ar;
6 }
7 void main() {
8     int i, *a = makeArray(N);
9     for(i=0; i<N; i++)
10         * (a+i) = i;
11     printf("a[] =
12         { %i, %i, %i } ", a[0], a[1], a[2]);
13     free(a);
14 }
```

- (A) 4, 12
- (B) 5, 12
- (C) 4, 10
- (D) 5, 10

Agenda

- C Memory Layout
 - Stack, Static Data, and Code
- Administrivia
- Dynamic Memory Allocation
 - Heap
- Common Memory Problems
- C Wrap-up: Linked List Example

Know Your Memory Errors

(Definitions taken from <http://www.hyperdictionary.com>)

- Segmentation Fault ← More common in 61C

“An error in which a running Unix program **attempts to access memory not allocated to it** and terminates with a segmentation violation error and usually a core dump.”
- Bus Error ← Less common in 61C

“A fatal failure in the execution of a machine language instruction resulting from the processor detecting an anomalous condition on its bus. Such conditions include **invalid address alignment (accessing a multi-byte number at an odd address)**, accessing a physical address that does not correspond to any device, or some other device-specific hardware error.”

Common Memory Problems

- 1) Using uninitialized values
- 2) Using memory that you don't own
 - Using NULL or garbage data as a pointer
 - De-allocated stack or heap variable
 - Out of bounds reference to stack or heap array
- 3) Using memory you haven't allocated
- 4) Freeing invalid memory
- 5) Memory leaks

Using Uninitialized Values

- What is wrong with this code?

```
void foo(int *p) {  
    int j;  
    *p = j; ← j is uninitialized (garbage),  
}  
                                         copied into *p
```

```
void bar() {  
    int i=10;  
    foo(&i);  
    printf("i = %d\n", i); ← Using i which now  
}  
                                         contains garbage
```

Using Memory You Don't Own (1)

- What is wrong with this code?

```
typedef struct node {  
    struct node* next;  
    int val;  
} Node;
```

```
int findLastNodeValue(Node* head) {  
    while (head->next != NULL)  
        head = head->next;  
    return head->val;  
}
```

What if `head`
is `NULL`?



No warnings!
Just Seg Fault
that needs finding!

Using Memory You Don't Own (2)

- What is wrong with this code?

```
char *append(const char* s1, const char *s2) {  
    const int MAXSIZE = 128; ← Local array appears  
    char result[MAXSIZE];          on Stack  
    int i=0, j=0;  
    for (j=0; i<MAXSIZE-1 && j<strlen(s1); i++,  
j++)  
        result[i] = s1[j];  
    for (j=0; i<MAXSIZE-1 && j<strlen(s2); i++,  
j++)  
        result[i] = s2[j]; ← Pointer to Stack (array)  
    result[++i] = '\0';           no longer valid once  
    return result;                function returns  
}
```

Using Memory You Don't Own (3)

- What is wrong with this code?

```
typedef struct {  
    char *name;  
    int age;  
} Profile;
```

```
Profile *person = (Profile *)malloc(sizeof(Profile));  
char *name = getName();  
person->name = malloc(sizeof(char) * strlen(name));  
strcpy(person->name, name);  
... // Do stuff (that isn't buggy)  
free(person);  
free(person->name);
```

Did not allocate space for the null terminator!
Want `(strlen(name)+1)` here.



Accessing memory after you've freed it.
These statements should be switched.

Using Memory You Haven't Allocated

- What is wrong with this code?

```
void StringManipulate() {  
    const char *name = "Safety Critical";  
    char *str = malloc(10);  
    strncpy(str, name, 10);  
    str[10] = '\0';           ← Write beyond array bounds  
    printf("%s\n", str);    ← Read beyond array bounds  
}
```

Using Memory You Haven't Allocated

- What is wrong with this code?

```
char buffer[1024]; /* global */  
int main(int argc, char *argv[]) {  
    strcpy(buffer, argv[1]);  
    ...  
}
```

What if more than
a kibi characters?

This is called **BUFFER OVERRUN or BUFFER OVERFLOW** and is a security flaw!!!

Freeing Invalid Memory

- What is wrong with this code?

```
void FreeMemX () {  
    int fnh = 0;  
    free (&fnh) ; ← 1) Free of a Stack variable  
}
```

```
void FreeMemY () {  
    int *fum = malloc (4*sizeof(int)) ;  
    free (fum+1) ; ← 2) Free of middle of block  
    free (fum) ;  
    free (fum) ; ← 3) Free of already freed block  
}
```

Memory Leaks

- What is wrong with this code?

```
int *pi;  
  
void foo() {  
    pi = (int*)malloc(8*sizeof(int));  
    ...  
    free(pi);      Overrode old pointer!  
}  
  
void main() {  
    pi = (int*)malloc(4*sizeof(int));  
    foo();         foo() leaks memory  
}
```

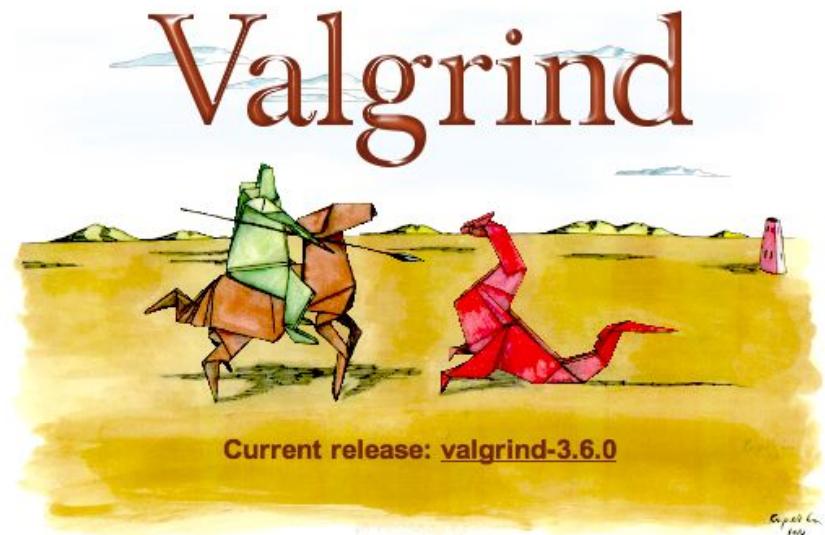
Memory Leaks

- **Rule of Thumb:** More mallocs than frees probably indicates a memory leak
- Potential memory leak: Changing pointer – do you still have copy to use with free later?

```
plk = (int *)malloc(2*sizeof(int));  
...  
plk++; ← Typically happens through incrementation  
or reassignment
```

Debugging Tools

- Runtime analysis tools for finding memory errors
 - Dynamic analysis tool:
Collects information on memory management while program runs
 - No tool is guaranteed to find ALL memory bugs;
this is a very challenging programming language research problem
- You will be introduced to Valgrind in Lab 3



<http://valgrind.org>

Get To Know Your Staff



Jinglin



Reese

Favorite Class	Music 128	CS61B
Favorite Place	Marina	Grizzly Peak
Favorite Restaurant	Spoon	Joshu-Ya
Favorite Club/Group	n/a	Cal Cycling

Agenda

- C Memory Layout
 - Stack, Static Data, and Code
- Administrivia
- Dynamic Memory Allocation
 - Heap
- Common Memory Problems
- C Wrap-up: Linked List Example

Linked List Example

- We want to generate a **linked list of strings**
 - This example uses structs, pointers, malloc (), and free ()
- Create a structure for nodes of the list:

```
struct Node {  
    char *value;  
    struct Node *next;  
};
```

The link of
the linked list

Simplify Code with `typedef`

- It gets annoying to type out `struct ListNode` over and over again
 - Define new variable type for struct:

Method 1: Method 2:

```
struct Node {  
    char *value;  
    struct Node *next;  
};  
typedef struct Node ListNode;
```

```
typedef struct Node {  
    char *value;  
    struct Node *next;  
} ListNode;
```

- Can further rename pointers:

```
typedef ListNode * List; → List myLinkedList;  
typedef char * String; → String value;
```

Adding a Node to the List

- Want functionality as shown:

```
String s1 = "start", s2 = "middle";
String s3 = "end";
List theList = NULL;
theList = addNode(s3, theList);
theList = addNode(s2, theList);
theList = addNode(s1, theList);
```

In what part of
memory are
these stored?

Must be able to
handle a
NULL input

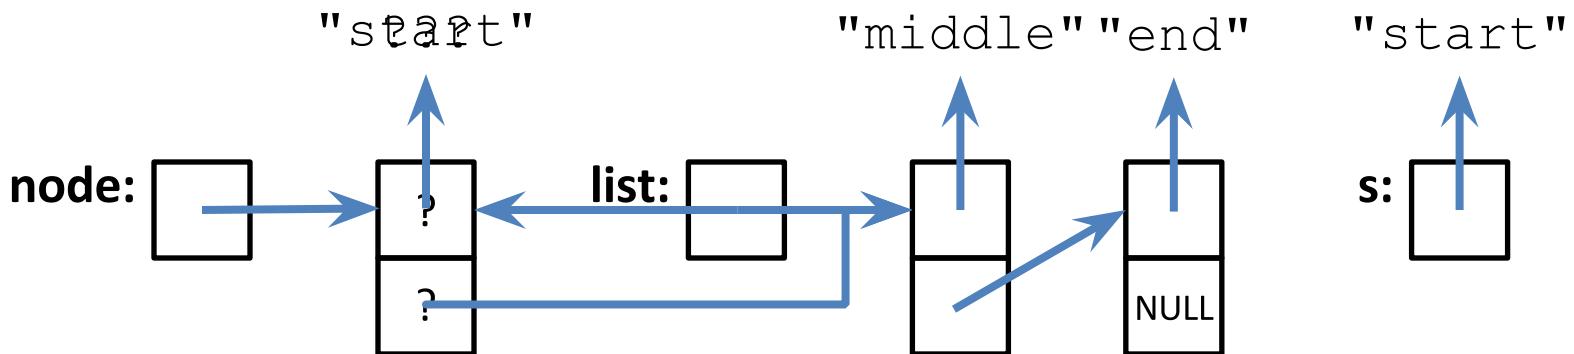
If you're more familiar with Lisp/Scheme,
you could name this function cons instead.

Adding a Node to the List

- Let's examine the 3rd call ("start"):

```
List addNode(String s, List list) {  
    List node = (List) malloc(sizeof(NodeStruct));  
    node->value = (String) malloc(strlen(s) + 1);  
    strcpy(node->value, s);  
    node->next = list;  
    return node;  
}
```

Don't forget this for
the null terminator!

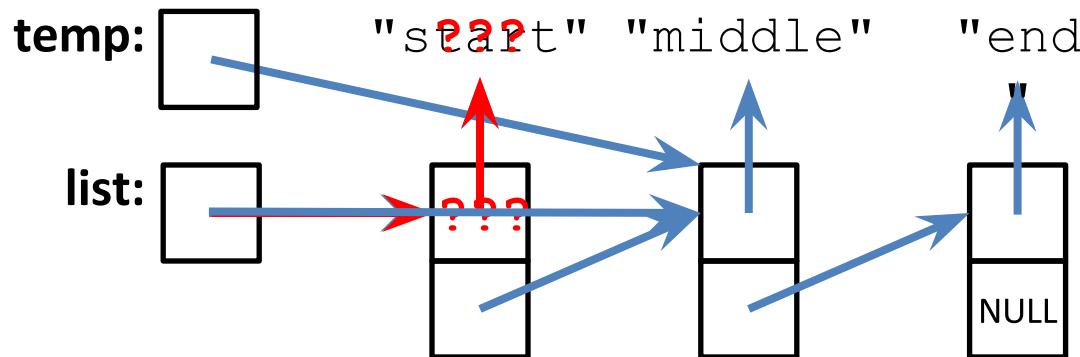


Removing a Node from the List

- Delete/free the first node ("start"):

```
List deleteNode(List list) {  
    List temp = list->next;  
    free(list->value);  
    free(list);  
    return temp;  
}
```

What happens if you do
these in the wrong order?



Additional Functionality

- How might you implement the following:
 - Append node to end of a list
 - Delete/free an entire list
 - Join two lists together
 - Reorder a list alphabetically (sort)

Summary

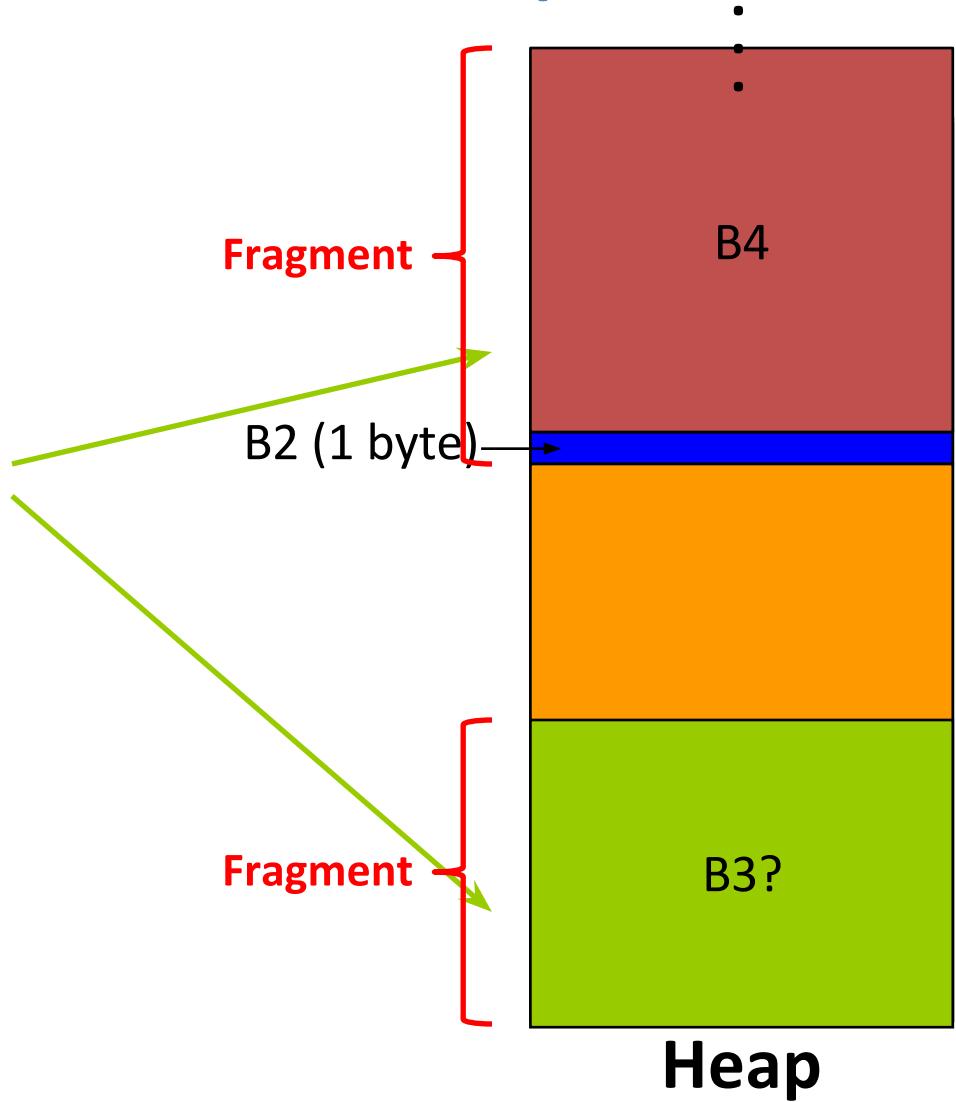
- C Memory Layout
 - **Static Data:** globals and string literals
 - **Code:** copy of machine code
 - **Stack:** local variables (grows & shrinks in LIFO manner)
 - **Heap:** dynamic storage using `malloc` and `free`
The source of most memory bugs!
- Common Memory Problems
- Last C Lecture!

Memory Management

- Many calls to `malloc()` and `free()` with many different size blocks – where are they placed?
- Want system to be fast with minimal memory overhead
 - Versus automatic garbage collection of Java
- Want to avoid *fragmentation*, the tendency of free space on the heap to get separated into small chunks

Fragmentation Example

- 1) Block 1: malloc(100)
- 2) Block 2: malloc(1)
- 3) Block 1: free(B1)
- 4) Block 3: malloc(50)
 - What if malloc(101)?
- 5) Block 4: malloc(60)



Basic Allocation Strategy: K&R

- Section 8.7 offers an implementation of memory management (linked list of free blocks)
 - If you can decipher the code, you're well-versed in C!
- This is just one of many possible memory management algorithms
 - Just to give you a taste
 - No single best approach for every application

K&R Implementation

- Each block holds its own **size** and **pointer to next block**
- `free()` adds block to the list, combines with adjacent free blocks
- `malloc()` searches free list for block large enough to meet request
 - If multiple blocks fit request, which one do we use?

Choosing a Block in malloc()

- **Best-fit:** Choose smallest block that fits request
 - Tries to limit wasted fragmentation space, but takes more time and leaves lots of small blocks
- **First-fit:** Choose first block that is large enough (always starts from beginning)
 - Fast but tends to concentrate small blocks at beginning
- **Next-fit:** Like first-fit, but resume search from where we last left off
 - Fast and does not concentrate small blocks at front